

# Reinforcement Learning for Highway Driving

Marcelo García Escalante  
mrgaresc@stanford.edu

Saba Bokredenghel  
sababo@stanford.edu

May 5, 2023

## 1 Problem Statement and Task Definition

Autonomous driving is a rapidly growing field of research. The goal of this project is to develop a reinforcement learning agent that can drive a car on a highway. The agent will be trained in a simulated environment and will be evaluated on its ability to drive safely (with out collisions) and efficiently (at the highest speed allowed). The agent will be trained to drive in a highway with multiple lanes. The agent will be evaluated on its ability to drive safely and efficiently.

The environment that we will use is built on top of OpenAI Gym.[1] The environment is called “highway-v0”[2] We will use this environment to train and evaluate the agent using reinforcement learning algorithms we have learned in class such as Value Iteration or Q-learning, compare their outcomes and discuss the results.

## 2 Environment Description

### 2.1 Input (State Space)

The input to the agent will be the state of the environment. The observation space is a continuous space 5x5 2D-array: `-inf, inf, (5, 5), float32`

The state space captures the ego-vehicle’s information plus the 4 closest vehicles in the highway, where the meaning of the features of the state are the following:

- **presence:** 1.0 if a vehicle is present, 0.0 otherwise.
- **x:** World offset of ego vehicle or offset to ego vehicle on the x axis.
- **y:** World offset of ego vehicle or offset to ego vehicle on the y axis.
- **vx:** Velocity on the x axis of vehicle.
- **vy:** Velocity on the y axis of vehicle.

**Note:** the coordinates are relative to the ego-vehicle, except for the ego-vehicle which stays absolute. The world frame is at the top left-corner of the highway. Examples of the state space are shown in appendix A

### 2.2 Output (Action Space)

The action space is a discrete space with 5 possible actions:

Action index	0	1	2	3	4
Action name	lane change left	idle	lane change right	accelerate	decelerate

## 3 Evaluation metric

The evaluation metric will be the average reward per episode. The reward function depends on 4 factors by default, but it can be changed.

- **Collision:** 0 if there is a collision, 1 otherwise.
- **Speed:** 0 to 1 depending on the speed of the ego-vehicle. 1 represents the maximum allowed speed
- **Close to right lane:** 0 to 1 depending on the distance to the right lane.
- **On road reward:** 0 or 1 depending on whether the ego-vehicle is on the road or not.

## 4 Baseline and Oracle

The baseline will be a random agent that selects an action uniformly at random. The oracle will be a human driver driving along the highway for 30 seconds as this is the maximum time allowed for the agent to drive in the environment per episode.

- 5 Related works
- 6 Methodology
- 7 Description of the challenges

## References

- [1] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. OpenAI Gym, 2016.
- [2] Edouard Leurent. An environment for autonomous driving decision-making. <https://github.com/eleurent/highway-env>, 2018.

## A Appendix: Raw state space samples

### Sample 1

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.1200	1.0	0.0
vehicle 1	1.0	0.1900	0.0700	-0.3700	-0.1800
vehicle 2	1.0	0.4300	0.1100	-0.2000	0.1400
vehicle 3	1.0	0.6400	0.0000	-0.2000	0.0000
vehicle 4	1.0	0.8800	0.0800	-0.1200	0.0000

### Sample 2

Table 1: Sample 2

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.0800	1.0	0.0
vehicle 1	1.0	0.2100	-0.0400	-0.1800	0.0000
vehicle 2	1.0	0.2400	-0.0800	-0.4800	0.0000
vehicle 3	1.0	0.5700	-0.0800	-0.3400	0.0000
vehicle 4	1.0	0.9500	-0.0800	-0.1600	0.0000

### Sample 3

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.0500	0.6200	-0.1400
vehicle 1	1.0	0.0100	-0.0300	-0.0900	0.1400
vehicle 2	1.0	-0.2400	-0.0500	0.2200	0.1400
vehicle 3	1.0	0.2500	-0.0500	0.3500	0.1400
vehicle 4	1.0	0.8000	0.0300	0.3300	0.1400

### Sample 4

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.0	1.0	0.0
vehicle 1	1.0	0.0600	0.0800	-0.4200	0.0000
vehicle 2	1.0	0.2500	0.0000	-0.4600	0.0000
vehicle 3	1.0	0.4800	0.0400	-0.5000	0.0000
vehicle 4	1.0	0.7700	0.0800	-0.4300	0.0000

### Sample 5

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.0800	1.0	0.0
vehicle 1	1.0	0.1600	-0.0400	0.0700	0.0000
vehicle 2	1.0	0.2200	-0.0800	-0.2000	0.0000
vehicle 3	1.0	0.5700	-0.0800	-0.0600	0.0000
vehicle 4	1.0	0.9300	-0.0100	-0.0500	-0.3300

### Sample 6

Vehicle	presence	x	y	vx	vy
ego-vehicle	1.0	1.0	0.0800	1.0	0.0
vehicle 1	1.0	-0.1600	-0.0800	-0.0300	0.0000
vehicle 2	1.0	0.2600	0.0000	-0.0200	0.0000
vehicle 3	1.0	-0.3300	-0.0400	-0.0600	-0.0100
vehicle 4	1.0	0.6800	-0.0800	0.0200	0.0000